

CS 440: Introduction to Artificial Intelligence

Homework #2 Search Problems in AI

March 10, 2023

Professor Abdelsam Boularias

Jay Patwardhan 208001851

Alan Wu 208000574

Neel Shejwalkar 207004853

Problem 1

Tracing Operation A^* from Lugoj to Bucharest

By applying the A^* tree search algorithm on the graph using the straight line distance heuristic, we end up with the following sequence of nodes expanded in order, written in tuples in the form described in the assignment:

- | | |
|------------------------------|------|
| $(Lugoj, 244, 0, 244)$ | (1) |
| $(Mehadia, 311, 70, 241)$ | (2) |
| $(Lugoj, 384, 140, 244)$ | (3) |
| $(Drobeta, 387, 145, 242)$ | (4) |
| $(Craiova, 425, 265, 160)$ | (5) |
| $(Timisoara, 440, 111, 329)$ | (6) |
| $(Mehadia, 451, 210, 241)$ | (7) |
| $(Mehadia, 461, 220, 241)$ | (8) |
| $(Lugoj, 466, 222, 244)$ | (9) |
| $(Pitesti, 503, 403, 100)$ | (10) |
| $(Bucharest, 504, 504, 0)$ | (11) |

Problem 2

Consider a state space where the start state is number 1 and each state k has two successors: numbers $2k$ and $2k + 1$.

Part A

Suppose the goal state is 11. List the order in which states will be visited for breadthfirst search, depth-limited search with limit 3, and iterative deepening search.

Solution. *The list of states visited for each search algorithm is as follows:*

For Breadth-First Search: (1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11)

For Depth-Limited Search, limit 3: (1, 2, 4, 8, 9, 5, 10, 11)

For Iterative Deepening search:

$Limit = 0$ (1)

$Limit = 1$ (1, 2, 3)

$Limit = 2$ (1, 2, 4, 5, 3, 6, 7)

$Limit = 3$ (1, 2, 4, 8, 9, 5, 10, 11)

Part B

How well would bidirectional search work on this problem? List the order in which states will be visited. What is the branching factor in each direction of the bidirectional search?

Solution. *The branching factor from the direction of the root node is 2, as every node has 2 children, and every parent node would have already been searched. The branching factor from the direction of the goal node is 3, as we have to account for its 2 children as well as the parent node.*

Asymptotically, bidirectional search is better than the other singular directional algorithms above (on the order of $O(2^{d/2} + 3^{d/2})$ versus $O(2^d)$) but when d is a small value like in this problem, it is possible for it to not perform as well.

The explored states are: (1, 11, 2, 3, 5, 22, 23, 4, 5).

Problem 3

Correct vs. Not Correct

Part A

Breadth-first search is a special case of uniform-cost search.

Part B

Depth-first search is a special case of best-first tree search.

Part C

Uniform-cost search is a special case of A^* search.

Part D

Depth-first graph search is guaranteed to return an optimal solution

Part E

Breadth-first graph search is guaranteed to return an optimal solution.

Part F

Uniform-cost graph search is guaranteed to return an optimal solution.

Part G

A^* graph search is guaranteed to return an optimal solution if the heuristic is consistent.

Part H

A^* graph search is guaranteed to expand no more nodes than depth-first graph search if the heuristic is consistent.

Part I

A^* graph search is guaranteed to expand no more nodes than uniform-cost graph search if the heuristic is consistent.

Problem 4

Iterative deepening is sometimes used as an alternative to breadth first search. Give one advantage of iterative deepening over BFS, and give one disadvantage of iterative deepening as compared with BFS. Be concise and specific.

An advantage that IDS has over BFS is its efficient use of space. BFS uses $O(b^d)$ space as it must store at worst every single node in the tree of a depth shallower than the goal node. In comparison, IDS is identical to Depth-First search during each one of its iterations. This means that IDS has a space complexity of $O(bd)$ or even $O(d)$ depending on the implementation, far better than the exponential space required for BFS.

IDS has the disadvantage of being much less efficient with time than BFS. While they are both similar asymptotically, with a complexity of $O(b^d)$, IDS needs to search a tree of depth i for every single iteration $0 \leq i \leq d$, whereas BFS needs to only search the tree of depth d once before finding the goal node. This means the runtime of IDS has a looser upper bound of $O(b^d + b^{d-1} + \dots + b + 1)$, which is far more expensive practically.

Problem 5

Prove that if a heuristic is consistent, then it must be admissible. Construct an example of an admissible heuristic that is not consistent. (Hint: You can draw a small graph of 3 nodes and write arbitrary cost and heuristic values so that the heuristic is admissible but not consistent.)

We will prove this claim using induction over all nodes.

Base case: Let n_0 be the goal node, and n_1 be a node one step away. By the definition of a consistent heuristic,

$$h(n_1) \leq c(n_1, a, n_0) + h(n_0)$$

$$h(n_1) \leq c(n_1, a, n_0) + 0$$

$$h(n_1) \leq h^*(n_1)$$

as the heuristic is defined to be 0 at the goal node, and the cost of the path from n_1 to n_0 is exactly the true cost of the path to the goal, $h^*(n_1)$. So, h demonstrates admissible behavior for the base case.

Inductive hypothesis: Let the condition $h(n_k) \leq h^*(n_k)$ hold for some node n_k that is k steps away from the goal node.

Inductive step: We'll now look at a node n_{k+1} that is $k+1$ steps away from the goal node. By the definition of consistency,

$$h(n_{k+1}) \leq c(n_{k+1}, a, n_k) + h(n_k)$$

By the inductive hypothesis,

$$h(n_{k+1}) \leq c(n_{k+1}, a, n_k) + h^*(n_k)$$

Now notice that $c(n_{k+1}, a, n_k) + h^*(n_k)$ is exactly the true distance from the node n_{k+1} . So, then

$$h(n_{k+1}) \leq h^*(n_{k+1})$$

And so the heuristic is admissible for all nodes. □

Problem 6

In a Constraint Satisfaction Problem search, explain why it is a good heuristic to choose the variable that is most constrained but the value that is least constraining.

Problem 7

Consider the following game tree, where the first move is made by the MAX player and the second move is made by the MIN player.

Part A


What is the best move for the MAX player using the minimax procedure?

The MAX player here has two moves, to go to state B (we'll call this action a_1) or go to state C (a_2). If MAX chooses a_1 , then MAX will end with a utility of 3. If MAX choose a_2 , it ends with a utility of 4 as the game progresses towards state H. We can check this result directly using the definition of the MINIMAX function:

$$\begin{aligned} \text{MINIMAX}(\text{root}) &= \max(\min(3, 5), \min(\max(\min(0, 7), 5), \max(7, 8), 4)) \\ &= \max(3, \min(5, 8, 4)) \\ &= \max(3, 4) = 4 \end{aligned}$$


Part B

Perform a left-to-right (left branch first, then right branch) alpha-beta pruning on the tree. That is, draw only the parts of the tree that are visited and don't draw branches that are cut off (no need to show the alpha or beta values).



Part C

Do the same thing as in the previous question, but with a right-to-left ordering of the actions. Discuss why different pruning occurs.



The differences in pruning arise because we are visiting the nodes in a different order. When going from left to right, the left node is checked with other states to see if pruning is a possibility, and if so, the right subtree is pruned. Conversely, going from right to left means that the right node is checked first and left subtrees are pruned. This difference in order is what changes the result. For example, when pruning left to right, we visited node K before node L. Seeing as the utility of node K was 7, we pruned off the right subtree of G (as the MIN player would rather go to node F regardless of what node K was). Similarly, when going from right to left, the left subtree of G was pruned (as the MIN player would have opted for node H).

Problem 8

Which of the following are admissible, given admissible heuristics h_1, h_2 ? Which of the following are consistent, given consistent heuristics h_1, h_2 ? Justify your ans

Part A

$$h(n) = \min(h_1(n), h_2(n))$$

Part B

$$h(n) = \omega h_1(n) + (1 - \omega)h_2(n), \text{ where } 0 \leq \omega \leq 1.$$

Part C

$$h(n) = \max(h_1(n), h_2(n))$$

Problem 9

Simulated annealing is an extension of hill climbing, which uses randomness to avoid getting stuck in local maxima and plateaux.

Part A

For what types of problems will hill climbing work better than simulated annealing? In other words, when is the random part of simulated annealing not necessary?

Part B

For what types of problems will randomly guessing the state work just as well as simulated annealing? In other words, when is the hill-climbing part of simulated annealing not necessary?

Part C

Reasoning from your answers to parts (a) and (b) above, for what types of problems is simulated annealing a useful technique? In other terms, what assumptions about the shape of the value function are implicit in the design of simulated annealing?

Part D

As defined in your textbook, simulated annealing returns the current state when the end of the annealing schedule is reached and if the annealing schedule is slow enough. Given that we know the value (measure of goodness) of each state we visit, is there anything smarter we could do?

Part E

Simulated annealing requires a very small amount of memory, just enough to store two states: the current state and the proposed next state. Suppose we had enough memory to hold two million states. Propose a modification to simulated annealing that makes productive use of the additional memory. In particular, suggest something that will likely perform better than just running simulated annealing a million times consecutively with random restarts. [Note: There are multiple correct answers here.]

Part F

Gradient ascent search is prone to local optima just like hill climbing. Describe how you might adapt randomness in simulated annealing to gradient ascent search avoid trap of local maximum.